Detecting Volcanoes on Venus via Classification

Team Mamamia!:Di Ye (diye2)[1], Hao Wang (haow2)[1], and Hannah Pu (chpu2)[1]

[1]Department of Statistics, University of Illinois at Urbana-Champaign

LASSO Regression, logistic regression, neural network, random forest, and xgboost have been used in this project to identify whether each image in the Kaggle datasets contains volcanoes or not, and if there is any, how many volcanoes are there. The analysis of the images revealed that there are several volcanoes on Venus. Neural network xgboost both had good performance to identify whether the images have volcanoes in this project for image classification with an area under the curve of 0.9629 (neural network) and 0.9680 (xgboost), and F-1 score of 0.9747 (neural network) and 0.9628 (xgboost). The comparison between different models is also included.

Image Classification | LASSO Regression | Logistic Regression | Marginal Screening | Neural Network | Random Forest | Xgboost | Kmeans | Linear Discriminant Analysis

I. INTRODUCTION.
It is well known that Venus is the closest planet to Earth. However, the surface of Venus is obscured by layers of thick cloud cover. It was until NASA's Magellan spacecraft was launched off for a mission to Venus on May 4, 1989 that the first detailed information about the surface of Venus was revealed. The mission mapped surfaces of Venus and image recorded most 98% of Venus surface. Dr. King in his article *Volcanoes on Venus* has mentioned that "volcanic activity is the dominant process for shaping the landscape of Venus, with over 90% of the planet's surface being covered by lava flows and shield volcanoes". Based on the images of volcanoes on Venus, we aim to construct prediction model to identify whether an image contains volcanoes or not, and how many volcanoes are there. **Figure 1** is a simulated color image of the surface of Venus created by NASA using radar topography data acquired by the Magellan spacecraft.

Data Source Information. The data was downloaded from Kaggle, which is originally from NASA's Magellan spacecraft database. (https://www.kaggle.com/amantheroot/finding-volcanoes-on-venus/data)

Data Description. 9734 images were captured by the NASA's Magellan spacecraft in the 1990s and converted to pixels (110 x 110, from 0 to 255), where each image is one row of 12100 columns (all the 110 rows of 110 columns). Images can contain more than one volcanoes or maybe none. The 9000+ images are separated to four datasets (file names : *train_images*, *train_labels*, *test_images*, and *test_labels*):

Fig. 1. Volcanoes on Venus

Image Dataset (*train_images* and *test_images*)
Train_images : 7000 images as train data with 12100 variables;
Test_images : 2734 images as test data with 12100 variables; All the variables (V1 to V12100) correspond to the pixel image, 110 pixels * 110 pixels = 12100 pixels.

Label Dataset (*train_labels* and *test_labels*) A summary of the variables in both *train_labels* and *test_label* datasets is listed down below:
1. *Volcano?* : If in the image there exists at least one volcano, the label is 1 (Yes). Otherwise, the label is 0 (No). (If Volcano? equals 0, the following three categories would be "NaN"). 2. *Type* : 1 = definitely a volcano,2 = probably, 3 = possibly, 4 = only a pit is visible
3. *Radius* : Is the radius of the volcano in the center of the image, in pixels?
4. *Number Volcanoes* : The number of volcanoes in the image.

Literature Review. Scientific papers regarding images obtained throught the Magellan mission were published, including Berl etc in 1989. In the 1989 paper, machine learning approach was introduced to facilitate in recognizing volcanoes from the images obtained in the mission. It was mentioned that in image data, volcanoes could appear in any pixel location, therefore principal components analysis (PCA) and linear discriminant analysis (LDA) could be implemented in feature extraction. Various training algorithms such as quadratic classifiers, decision trees, linear discriminant analysis, and nearest neighbors were used and produced similar outcomes.

These venus datasets were uploaded in Kaggle and 11 kernels regarding this data analysis were produced, all using Python. The kernels included vivid data visualization and exploratory data. Most kernels focused on classification prediction on whether there are volcanoes or not. For this classification prediction, different methods have been used to construct prediction model, from simple models such as logistic regression, to advanced models such as Convolutional Neural Network (CNN) and VGG Neural Network for deep learning. Accuracy of models ranged from 84.1% to 97%.

Scientific Goal. For this project, we focus mainly on doing classification predition on whether each image has a volcano or not. We aim in constructing different classification models and choosing the best model to predict whether there exists a valcano in each image. Identifying valcanoe through IT technology would increase the efficency of space exploration and safty of the crews. In addition, we also constructed a few models to perform prediction on the number of volcanoes in the images.

II. EXPLORATORY DATA. The first 6 observations of *train_labels*

```
head(train_y)
```

```
#     Volcano. Type Radius Number.Volcanoes
# 1          1    3  17.46                1
# 2          0  NaN    NaN              NaN
# 3          0  NaN    NaN              NaN
# 4          0  NaN    NaN              NaN
# 5          0  NaN    NaN              NaN
# 6          0  NaN    NaN              NaN
```

The first 6 observations of *test_labels*

```
head(test_y)
```

```
#     Volcano. Type Radius Number.Volcanoes
# 1          0  NaN    NaN              NaN
# 2          0  NaN    NaN              NaN
# 3          1    1  17.00                1
# 4          0  NaN    NaN              NaN
# 5          1    3  15.13                1
# 6          0  NaN    NaN              NaN
```

After exploring the datasets, we found only labels have NaNs. The NaNs in the label dataset occurs only when there is no Volcanoe in the image observation, then `Type`, `Radius` and `Number.Volcanoes` would be NaNs as there are no volcanoes. We have set the all these NaNs to 0.

```
test_y[is.na(test_y$Type), ] <- 0
test_y[is.na(test_y$Radius), ] <- 0
test_y[is.na(test_y$Number.Volcanoes), ] <- 0
train_y[is.na(train_y$Type), ] <- 0
train_y[is.na(train_y$Radius), ] <- 0
train_y[is.na(train_y$Number.Volcanoes), ] <- 0
```

The first 6 observations of *train_labels* after setting NaNs to 0

```
head(train_y)
```

```
#     Volcano. Type Radius Number.Volcanoes
# 1          1    3  17.46                1
# 2          0    0   0.00                0
# 3          0    0   0.00                0
# 4          0    0   0.00                0
# 5          0    0   0.00                0
# 6          0    0   0.00                0
```

The first 6 observations of *test_labels* after setting NaNs to 0

```
head(test_y)
```

```
#     Volcano. Type Radius Number.Volcanoes
# 1          0    0   0.00                0
# 2          0    0   0.00                0
# 3          1    1  17.00                1
# 4          0    0   0.00                0
```

Fig. 2. Obs 10 Type: 1 Radius: 22.02 Number Volcanoes: 1

Fig. 3. Obs 39 Type: 2 Radius: 19.31 Number Volcanoes: 1

```
#  5          1     3   15.13                    1
#  6          0     0    0.00                    0
```

Data Visualization. 6 observations are picked to demonstrate how the images got labeled. **Figure 2 - 5** show how well the volcanoes can be seen from the images (*Type* : 1 = definitely a volcano, 2 = probably, 3 = possibly, 4 = only a pit is visible). **Figure 6** contains 2 volcanoes, and Figure 6 contains no volcano at all. We can tell from the images that a bright white dot indicates a potential volcano, while the white dot might not be clear enough to see in the image, which means that it is hard to identify whether there is a volcano or not.

Fig. 4. Obs 1 Type: 3 Radius: 17.46 Number of Volcanoes: 1

The following `volplot` function is written for plotting some sample images.

```
volplot <- function(data, obs){
  im <- as.numeric(data[obs,])
  m <- matrix(im, nrow = 110, byrow = TRUE)
  image(m, col = grey((0:255)/255))
}
```

An example code for plotting the 10th observation in the *train_images*

```
volplot(train_data, 10)
```


**Fig. 5.** Obs 30 Type: 4 Radius: 6.40 Number Volcanoes: 1

Fig. 6. Obs 289 Type: 1 Radius: 11.05 Number Volcanoes: 2


Fig. 7. Obs 2 No Volcano

Fig. 8. How many images have volcanoes?

Data Summary. To facilitate the analysis process and explore the data, we have summarized the data. The `ggplot` has been used to visualize the training set labels. **Figure 8** shows that only 1000 images in the training dataset have volcanoes. **Figure 9** shows that within 1000 images that have volcanoes, how clearly we can identify the volcanoes. **Figure 10** shows how many volcanoes there are in each image. The number of volcanoes ranges from 0 to 5.


**Fig. 9.** How clealy can we identify the volcanoes?

Fig. 10. How many volcanoes are there in each image?

III. METHODOLOGY & ANALYSIS. As the project data is image data, preprocessing was not neccessary to fit model. We studied in performing two types of prediction. The first type is to predict of whether there is volcanoe or not, we constructed simple prediction models and advanced models, which are discussed in detail later on. For evaluation metric for these models, as references has indicated that uncertainty exists in the identification of volcanoes in the images, we used AUC (Area under the ROC, receiver operating characteristic) and F1 score. AUC and F1 score provides overall performance measure of classifcation. the second type of prediction is to predict how many volcanoes there are, root mean square error (RMSE) was used for logistic, lasso, and xgboost and AUC was used for neural network. Models used for these two types of prediction are summarized as the following table.

Model	Predict if have volcano	Predict num of volcanoes
Neural Network	V	V
xgboost	V	V
Lasso	V	V
Random Forest	V	
Logistic (Marginal Screening)	V	V
Logistic (Lasso Selection)	V	V
Linear Discriminant Analysis	V	

A. Classification for if there is any volcano. Prior to using supervised learning methods, we applied kmeans clustering to the train dataset to see whether the pixels could accurately cluster the training dataset into 2 clusters. We used `kmeans` function with 2 centers and 20 random sets. The confusion matrix of the clusering

result are as following, it can be seen that by simply clustering the pixels, it is difficult to cluster the observations into 2 groups.

Cluster / Label	0	1
1	5810	999
2	190	1

Linear Dicriminant Analysis (LDA) and Quadratic Discriminant Analysis (QDA). Linear Dicriminant Analysis was applied by using `MASS` package. Prediction using LDA analysis produced adequate results, however, it output warning of 'variables are collinear'. Implementing QDA analysis produced error of 'some group is too small for qda'. The warning and error produced may be that our parameter (pixel) number is too large. Therefore to solve this problem, we try to use other algorithms to fit the dataset.

Lasso Regression. Lasso regression was applied as it can perform variable selection, which is suitable in the project situation : having more parameters than observations. To apply Lasso Regression, `glmnet` package was implemented. We first selected reasonable interval for tuning parameter λ by examing cross-validation error. After that, we did the model selection by using glmnet and use the "lambda.min" as final tuning parameter. By using the cross-validation error data, we identified the interval, (-6.5, -5) selected 470 important pixels. The prediction using those pixels yielded a good start, a high prediction accuracy.

Logistic Regression. Before performing logistic regression, we used two variable selection methods to reduce dimension. 400 pixels were selected by using Marginal Screening and 470 pixels were selected by using above Lasso Regression. Then, we applied logistic regressions on those 2 sets of selected pixels repectively. We then recorded their F1 score and AUC for classification, and they showed that Lasso selected pixels performed better than Marginal Screening method.

```
# code for marginal screening
cl <- makeCluster(3)
registerDoParallel(cl)
pre = Sys.time()
vol_results <- foreach(j = 1:ncol(train.x),
    .packages = "lme4", .combine = "rbind") %dopar%
    {
        pix = scale(train.x[, j])
        fit.glm <- summary(glm(as.factor(train.y) ~
            pix, family = "binomial"))
        fit.glm$coefficients[2,
```

```
            4]
    }
Sys.time() - pre
stopCluster(cl)
# save(vol_results, file =
# 'vol_result.RData')
rownames(vol_results) = colnames(train.x)
colnames(vol_results) = "glm.p value"
totalpix = 400
sortedresults = as.matrix(vol_results[order(vol_r
    ])
usepix = names(sortedresults[1:totalpix,
    ])
top_pix = match(usepix, rownames(vol_results))
```


Fig. 11. Neural Network to classify whether an image contains volcanoes

The loss function and accuracy:

```
#  $loss
#  [1] 0.3018637
#
#  $acc
#  [1] 0.9568398
```

```
# code for logistic regression:
logmodel = glm(as.factor(volcano) ~
    ., data = training, family = "binomial")
train.error = (predict(logmodel,
    data.frame(sel.pix), type = "response") >
    0.5)
log_pred = (predict(logmodel, data.frame(test.pipx),
    type = "response") > 0.5)
table(log_pred, test.y)
```

Neural Network. Convolutional neural networks (CNNs) have been applied widely in image and video recognition and classification. CNNs are made up of neurons with learnable weights and bias. Therefore, CNNs were used in the project for seeking better classfication results. Some pooling layers were added into the models to progressively reduce the spatial size and the amount of parameters and computation in the network. The raw data *train_images* and *test_images* was converted into numeric variables and reshaped into 3D array with dimension 7000 × 110 × 110 and 2734 × 110 × 110 respectively. The pixels were shrinked from 0-255 to 0-1 by dividing each pixel by 255. The target label (Volcano?) was converted into categorical variable for classification. The neural network models yielded efficient and satisfying classification results. The loss function is relatively low, indicating the models are good.

Figure 11 shows the result of the first neural network model which is designed to classify whether an image contains a volcano. The model used 20% of the training data as validation data.

Random Forest. Random Forest was implemented using the randomForest package. Originally, n.tree was set to 200 and all other parameters were set as default. However, due to our large parameter and observation number, this took over than 10 hours to run. As it took too much time, it was terminated before it finished running. The final model used in this project was by setting n.tree to 200, mtry to 80 and nodesize to 70. Variable importance were recorded to compare with variable selection in other methods. Variable importance from random forest are shown in **Figure 13**. The yellow represents the important ones while black means those pixels are less meaningful for our prediction.

```
# code for random forest
rfmodel <- randomForest(train.x,
    y = train.y[, 1], ntree = 200,
    mtry = 80, nodesize = 70)
```

```
# code for random forest
# variable important heatmap
library(RColorBrewer)
myimport <- rfmodel$importance
my_matrix <- matrix(myimport[,
    4], 110, 110, byrow = T)
pal <- colorRampPalette(brewer.pal(11,
    "RdYlGn"))(100)
mycol <- c("black", "yellow")
pal <- colorRampPalette(mycol)(100)
```

```
heatmap(my_matrix, Rowv = NA, Colv = NA,
    revC = T, col = pal)
```


Fig. 12. The Random Forest variable importance heatmap

Xgboost. As xgboost is nowadays a popular algorithm used in machine learning and Kaggle competitions, this algorithm was also used in this project to see its performance. Xgboost was implemented using the `xgboost` package. Tuning was performed by using `xgb.cv` function with nfold = 5, parameters that were tuned are eta (learning rate), max_depth (maximum depth of a tree), min_child_weight (minimum sum of instance weight needed in a child), subsample (subsample ratio of the training instance), colsample_bytree (subsample ratio of columns when constructing each tree). The tuning process took over 12 hours and was terminated before it finished all the combinations. The parameter combination with the lowest error was used. From the tuning process, parameterts used for this model are set as following: eta = 0.1, max_depth = 5, min_child_weight =5, subsample = 0.65, colsample_bytree = 0.8, nrounds = 200, objective = 'binary:logistic'. To model how many volcanoes there are, objective = 'count:possion' was used with same parameters as before.

```
#code for xgboost
params <- list(
  eta = 0.1,
  max_depth = 5, #7
  min_child_weight =5,
  subsample = 0.65,
  colsample_bytree = 0.8,
  silent = 1
```

```
)
xgb.fit.final <- xgboost(
  params = params,
  data = as.matrix(train.x),
  label = train.y,
  nrounds = 200,
  objective = 'binary:logistic',
  verbose = 0
)
```

Our model prediction results for classfication are summarized in table below.

Table 1: Test dataset prediction result on whether there is volcanoe or not

Model	AUC	F1 Score
Neural Network	0.9629	0.9747
xgboost	0.9680	0.9628
Lasso	0.9347	0.9574
Random Forest	0.9367	0.9399
Logistic (Marginal Screening)	0.8429	0.9386
Logistic (Lasso Selection)	0.8672	0.9527
Linear Discriminant Analysis	0.8893	0.9381

Xgboost and neural network both had good evaluation performance, their confusion matrix are as below:

Table 2: Classification results using neural network for identifying volcanoes in the images

	No	Yes
No	2270	88
Yes	30	346

Table 3: Classification results using xgoost for identifying volcanoes in the images

	No	Yes
No	2280	156
Yes	20	278

B. Poisson Prediction for how many volcanoes are there.

Lasso Regression:. For predicting the count of volcanoes, we also performed Lasso regression but changed the `family` to "poisson", which allows us to predict the count instead binary outcome. Similarly, we used cross-validation error to find the reasonable interval for lambda, (-5.5,-4) and used the lambda.min to predict. This approach gave us 114 important variables and RMSE 0.4216523.

```
# code for lasso
cv.out <- cv.glmnet(x = train.x,
    y = train.y, alpha = 1, family = "binomial")
lam.seq <- exp(seq(-6.5, -5, length = 100))
lasmodel <- glmnet(x = train.x,
```

```
     y = train.y, alpha = 1, family = "binomial",
     lambda = lam.seq)
```

Logistic Regression:. As we did in classification section, we used both Marginal Screenign and Lasso to select important pixels. In this case, Marginal Screening gave us 250 important pixels based on best prediction accuracy and RMSE 0.5072927; Lasso selected 114 pixels gave us RMSE 0.6798163.

Xgboost:. In performing prediction with xgboost, the tuning parameters were directly adopted from the xgboost model in the previous section.

Neural network:. We implemented similar algorithm for Neural network as in the previous section.

```
# Neural Network for
# number_volcano #

# convert the variable in to
# categorical variable
train_y_number <- to_categorical(train_y$Number.Volcanoes,
    num_classes = 6)
test_y_number <- to_categorical(test_y$Number.Vol
    num_classes = 6)

# model using convolutional
# layers and max pooling
model <- keras_model_sequential() %>%
    layer_conv_2d(filters = 8,
        kernel_size = c(3, 3),
        activation = "relu", input_shape = c(110,
            110, 1)) %>% layer_max_pooling_2d(poo
    2)) %>% layer_conv_2d(filters = 16,
    kernel_size = c(3, 3), activation = "relu") %>%
    layer_max_pooling_2d(pool_size = c(2,
        2)) %>% layer_conv_2d(filters = 16,
    kernel_size = c(3, 3), activation = "relu") %>%
    layer_flatten() %>% layer_dense(units = 6,
    activation = "softmax")

# compile model using loss
# function 'categorical
# crossentropy' and metrics
# 'accuracy'
model %>% compile(optimizer = "adam",
    loss = "categorical_crossentropy",
    metrics = c("accuracy"))

# use 20% of the train data as
# validation data
model_ret <- model %>% fit(train_images,
    train_y_number, epochs = 30,
```

```
    batch_size = 32, validation_split = 0.2)
plot(model_ret)
# use the trained model to
# classify test set
results <- model %>% evaluate(test_images,
    test_y_number)
# classification probabilities
pred_results_number <- model %>%
    predict(test_images)
# classification results
pred_results_number_class <- model %>%
    predict_classes(test_images)
# accuracy and loss function
# results
acc_loss_number <- model %>% evaluate(test_images,
    test_y_number)
```

Figure 12 shows the result of the second neural network model which is designed to classify whether an image contains at least one volcano. The model, similar to the first one, used 20% of the training data as validation data.


Fig. 13. Neural Network to classify how many volcanoes in each image

The loss function and accuracy:

```
#  $loss
#  [1] 0.554538
#
#  $acc
#  [1] 0.9297732
```

Our model prediction results for Poisson count prediction are summarized in table below. It can be seen that xgboost model has the lowest RMSE. As the neural network was a classification model, RMSE was not calculated. Instead, its confusion matrix are shown below. The number of volcanoes in images in the test set ranges from 0 to 2. However, in our classification,

the neural network identifies 12 images having 3 volcanoes. Other than that, the neural network yielded a good result.

Table 4: Test dataset prediction result on how many volcanoes there are in the images

Model	RMSE
Logistic	0.5073
xgboost	0.0613
Lasso	0.4613

Table 5: Classification results using neural network for identifying the number of volcanoes in the images

	0	1	2	3
0	2272	67	15	2
1	23	266	44	7
2	5	26	4	3

c. Variable selection methods.

Shrinking method (Lasso) By applying Lasso Regression, we were able to shrink highly correlated pixels to 0 and only had important ones left. With the tuning parameter λ's interval we selected by examing the cross validation result, we selected 470 important pixels for classification, 250 pixels for Poisson prediction, and apply those to lasso and logistic regression respectively.

Marginal Screening To apply Marginal Screening, we use for loop to model each pixels with outcome by using glm function, logistic regression. After sorting the p values from smallest to largest, we selected the top 400 pixels for classification and 250 for poisson prediction. The reason why we select top 400 and 250 pixels is that, after trying different number of top pixels, we found 400 pixels gave us highest prediction accuracy.

Variable selection methods comparison We originally only applied Marginal Screening to logistic regression. But after realizing Lasso helped us selected variables as well, we tried both methods on logistic regression.

By looking at the pixels' index, **figure 14** and **figure 16**, we found that the pixels selected by Marginal Screening focused between 4000-7000 for classification and between 5200 - 7000 for poisson prediction, which corresponds to the center of pictures; while pixels selected by lasso regression were more spread, as shown in **figure 15** and **figure 17**. Comparing these variable selection results to the variable importance in random forest, it is noted that trend of selected variables in Lasso corresponds to the variable importance in random forest.

By applying those two different selection methods, we found the accuracy of Lasso selection is higher for classification; while for Poisson prediction, Marginal Screening selection performs better.


**Fig. 14.** The Marginal Screening index for classification


Fig. 15. The Lasso index for classification

Marginal Screening index for Poisson prediction


Fig. 16. The Marginal Screening index for Poisson prediction

lasso index for Poisson prediction

Fig. 17. The Lasso index for Poisson prediction

Variable Importance in Random Forest

Fig. 18. Variable Importance in Random Forest

Hence, for classification, lasso's variable selection with more spread pixels made more sense because Marginal Screening's centered selection might ignore some important pixels at the edges. For Poisson prediction, we prefer the Marginal Screening since the smaller number of pixels selected by Lasso regression would omit some important information.

Overall, we could conclude the variable selection criteria for this project is how much information the index can tell us without overfitting model.

IV. CONCLUSION & DISCUSSION. Machine learning methods have long been popular in establishing models to faciliate humans in retreiving information from large datasets. One of the discussion topics in the venus image data, obtained through NASA's Magellan spacecraft, is to contruct models in facilitating the identification of whether there are volcanoes or not, how many there are, and the confidence people have in identifying they are volcanoes. For our project, we aim in establishing prediction models in predicting the existance of volcanoe in image data as well as establishing prediction models in predicting how many volcanoes there are. As the daataset is rather unbalanced, with 85% of our image data having no volcanoes, we use the evaluation metric of AUC and F1 score as our prediction metric.

For the prediction of "predicting whether there is volcanoe or not", our models show prediction results using AUC ranged from 0.8429 to 0.9680, using F1 score ranged from 0.9381 to 0.9747. When considering the best model chosen by either AUC or F1 score, it results in different models. For AUC, the best model is xgboost with AUC of 0.9680, whereas for F1 score the best model is neural network with F1 score of 0.9747. When comparing the confusion matrix of these two models, we notice that there are less false positive in xgboost where as there are less false negative in neural network.

For the prediction of "predicting how many volcanoes there are", our models show prediction results using RMSE ranged from 0.0613 to 0.5073. However the tuning parameterts of these prediction models were directly adopted from the model in the first prediction topic ("predicting whether there is volcano or not"). Lasso and logistic model performed poor results, compared to xgboost. As venus image data are unstructered and have noise, it is possible that prediction results would improve in these two models if dimension reduction techniques were applied. In future improvement, re-tuning these parameters would have improvement in predicting how many volcanoes there are.

Compared to structured data, the images used in this project are unstructured as volcanoes can occur in any pixels, this makes featuring difficult without performing principle component analysis. Models such as lasso, logistic, random forest, require more of pre-defined features. In this project, convolution neural network fits well in this kind of dataset as it considers local region of the datset.

It should be noted that, although the prediction result of Lasso regression is not as good as the advanced algorithms, this mode took significantly less running time. If to take into account of running time, Lasso regression is still a good choice.

The limit of the models in this project is that, due to large parameter and sample size. Tuning advanced models such as random forest, xgboost, and neural network takes a lot of time. The tuning done in this project could be more fined tuned in the future. When predicting the number of volcanoes, parameters in xgboost and neural network were directly used from the previous model, therefore these models could also be further tuned in the future.

In addition, dimension reduction methods such as single value decomposition was not performed. As image data are unstructered and image data of venus may include several noises, performing dimension reduction may have improvment in the prediction. As the dataset in this project was quite large, we were not able to use `svd` function to perform dimension reduction. For future improvement of this project, we could try faster packages such as `irlba` package to perform partial singular value decompositions and see if there is improvement in our model.

V. REFERENCES. Burl, M.C., Asker, L., Smyth, P. et al. Learning to Recognize Volcanoes on Venus. Machine Learning (1998) 30: 165. https://doi.org/10.1023/A: 1007400206189

Kumar, A., Finding Volcanoes On Venus. https://www.kaggle.com/amantheroot/finding-volcanoes-on-venus

King, H., Volcanoes on Venus. https://geology.com/stories/13/venus-volcanoes/